

NATURAL LANGUAGE PROCESSING

Hyperparameter Tuning in Pytorch

goorm

KAIST AI
Graduate School of AI



Contents

1. Types of Hyperparameters
2. Hyperparameter Optimization Steps
3. Further Hyperparameters
 - a. Accumulation Steps
 - b. Random Seed
 - c. Number of Evaluation
4. Wandb

Type of Hyperparameters

- Model-free hyperparameters
 - Learning rate
 - Batch size per gpu
 - Training epoch
 - Learning rate scheduler(Warm up steps, lambda, step size ...)
 - Optimizer (beta1, beta2 in Adam)
 - Weight initialization
 - Early stop strategy
 - Regularization
 - Dropout
 - Perturbation or noise for an input
 - ...

Type of Hyperparameters

- Model hyperparameters
 - Kernel size
 - number of layer
 - number of hidden units
 - number of embedding units
 - pooling
 - activation function

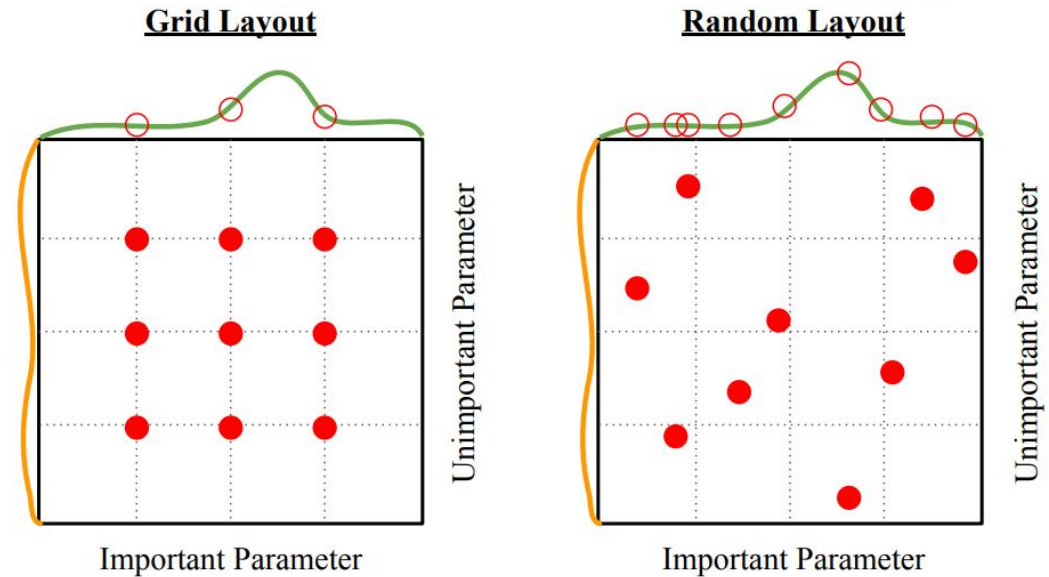
Type of Hyperparameters

- Model free & Model hyperparameters
 - Learning rate x Batch size per gpu x Training epoch x Learning rate scheduler (Warmup steps, lambda, step size...) x Optimizer (+beta1, beta2 in Adam) x Weight initialization x Early stop strategy x Regularization x Dropout x Kernel size x number of layer x number of hidden units x number of embedding units x pooling layer x activation function x....
 - [training time for a model](#)

Hyperparameter Optimization

- Hyperparameters Optimization

- Grid Search
- Random Search
- Bayesian



- (Maybe for researchers) Best practice for hyperparameters optimization

- ***Learning rate, Learning rate, Learning rate***
- ***Batch size***

Hyperparameter Optimization Steps

1. Check initial loss
2. Overfit a small sample
3. Find learning rate that makes loss go down
4. Coarse grid, train for 1~5 epochs
5. Refine grid, train longer than step 4
6. Look at loss curves

Hyperparameter Optimization Steps

I. **Check initial loss**

- a. Turn off learning rate scheduler and train a model

Hyperparameter Optimization Steps

1. Check initial loss
2. **Overfit a small sample**
 - a. Try to train to high training accuracy on a small sample of training data (~5-10 minibatches); fiddle with architecture, learning rate, weight initialization
 - b. Loss not going down? LR too low, bad initialization
 - c. Loss explodes to Inf or NaN? LR too high, bad initialization

Hyperparameter Optimization Steps

1. Check initial loss
2. Overfit a small sample
3. **Find learning rate that makes loss go down**
 - a. Use the architecture from the previous step, use all training data, turn on small weight decay, find a learning rate that makes the loss drop significantly within ~ 100 iterations
 - b. Good learning rates to try: $1e-1$, $1e-2$, $1e-3$, $1e-4$

Hyperparameter Optimization Steps

1. Check initial loss
 2. Overfit a small sample
 3. Find learning rate that makes loss go down
 4. **Coarse grid, train for 1~5 epochs**
 - a. Choose a few values of learning rate and weight decay around what worked from Step 3, train a few models for ~1-5 epochs.
 - b. Good weight decay to try: $1e-4$, $1e-5$, 0
- additional
 - Make possible batch sizes larger

Hyperparameter Optimization Steps




1. Check initial loss
2. Overfit a small sample
3. Find learning rate that makes loss go down
4. Coarse grid, train for 1~5 epochs
5. **Refine grid, train longer than step 4**
 - a. Pick best models from Step 4, train them for longer (~10-20 epochs) without learning rate decay

Hyperparameter Optimization Steps

1. Check initial loss
2. Overfit a small sample
3. Find learning rate that makes loss go down
4. Coarse grid, train for 1~5 epochs
5. Refine grid, train longer than step 4
6. **Look at loss curves**




Look at Learning Curves

I) Look at the training learning curves



Look at Learning Curves

2) Look at the training & validation learning curves



Hyperparameter Optimization

- Tips for Geeks who struggle for state of the art performance or want to beat competitors
 - Train your model with the largest batch size that memory allows in single gpu
 - Evaluate as many checkpoints as possible
 - Before the test, combine train set and validation set and train the model with the combined dataset
 - Do ensemble
 - The hyperparameters configuration of competitors who use similar model is a good starting point
 - Use MLOps tool (e.g., wandb)

Further Hyperparameters

- 1) Accumulation steps
- 2) Random seed
- 3) Number of evaluation

Further Hyperparameters

I) Accumulation steps

- 16 batch with no accumulation vs 4 batch with 4 accumulation step, which can result in better performance?

```
1 predictions = model(inputs)           # Forward pass
2 loss = loss_function(predictions, labels) # Compute loss function
3 loss.backward()                       # Backward pass
4 optimizer.step()                      # Optimizer step
5 predictions = model(inputs)           # Forward pass with new parameters
```

pytorch_training.py hosted with ❤ by GitHub

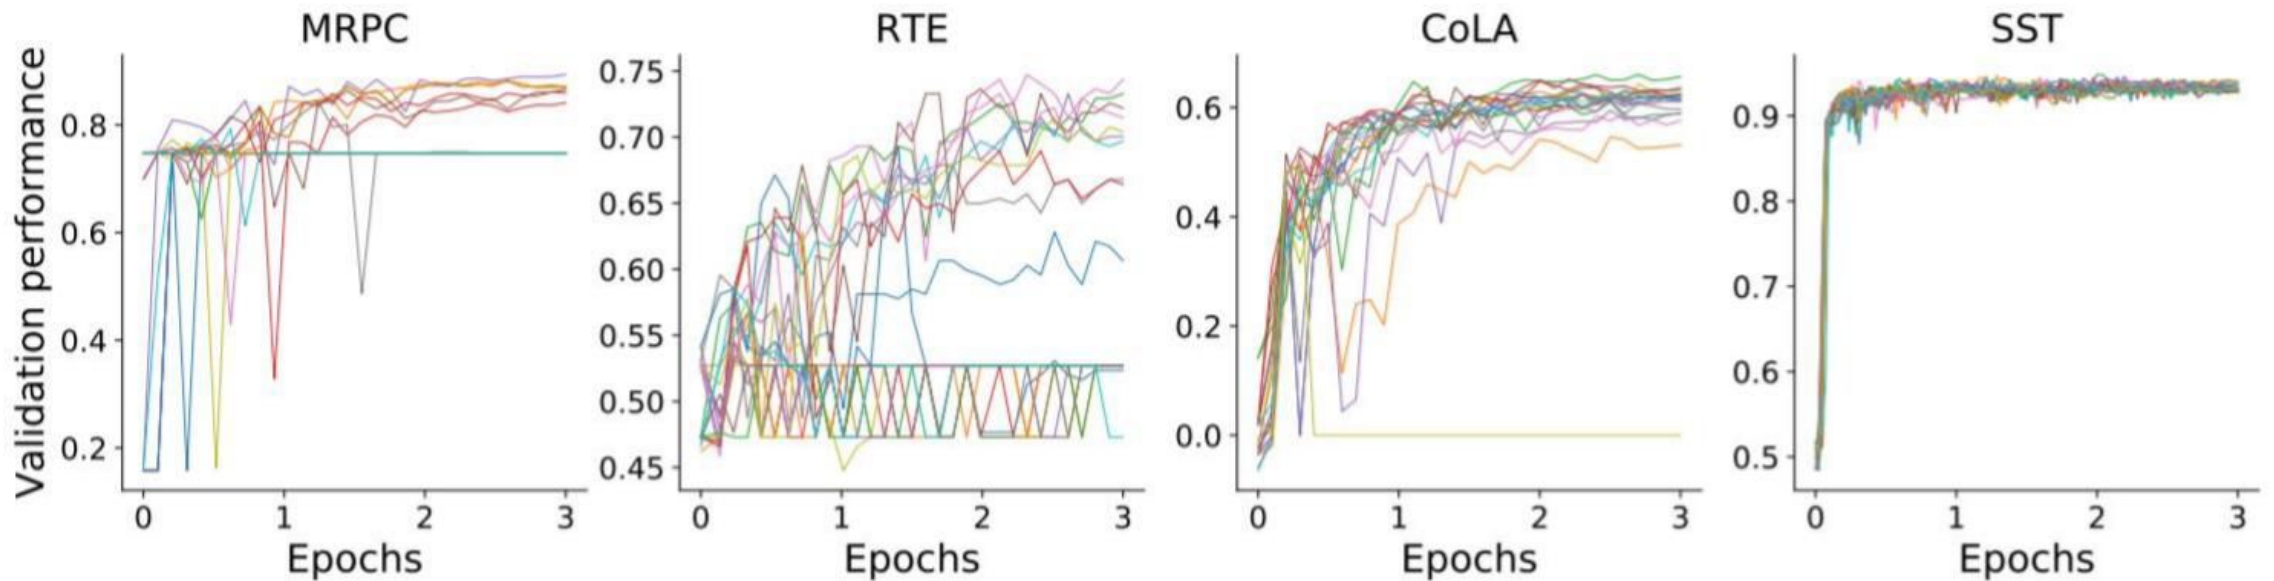
```
1 model.zero_grad()                     # Reset gradients tensors
2 for i, (inputs, labels) in enumerate(training_set):
3     predictions = model(inputs)         # Forward pass
4     loss = loss_function(predictions, labels) # Compute loss function
5     loss = loss / accumulation_steps    # Normalize our loss (if averaged)
6     loss.backward()                     # Backward pass
7     if (i+1) % accumulation_steps == 0: # Wait for several backward steps
8         optimizer.step()                # Now we can do an optimizer step
9         model.zero_grad()               # Reset gradients tensors
10        if (i+1) % evaluation_steps == 0: # Evaluate the model when we...
11            evaluate_model()              # ...have no gradients accumulated
```

gradient_accumulation.py hosted with ❤ by GitHub

Further Hyperparameters

2) Random Seed

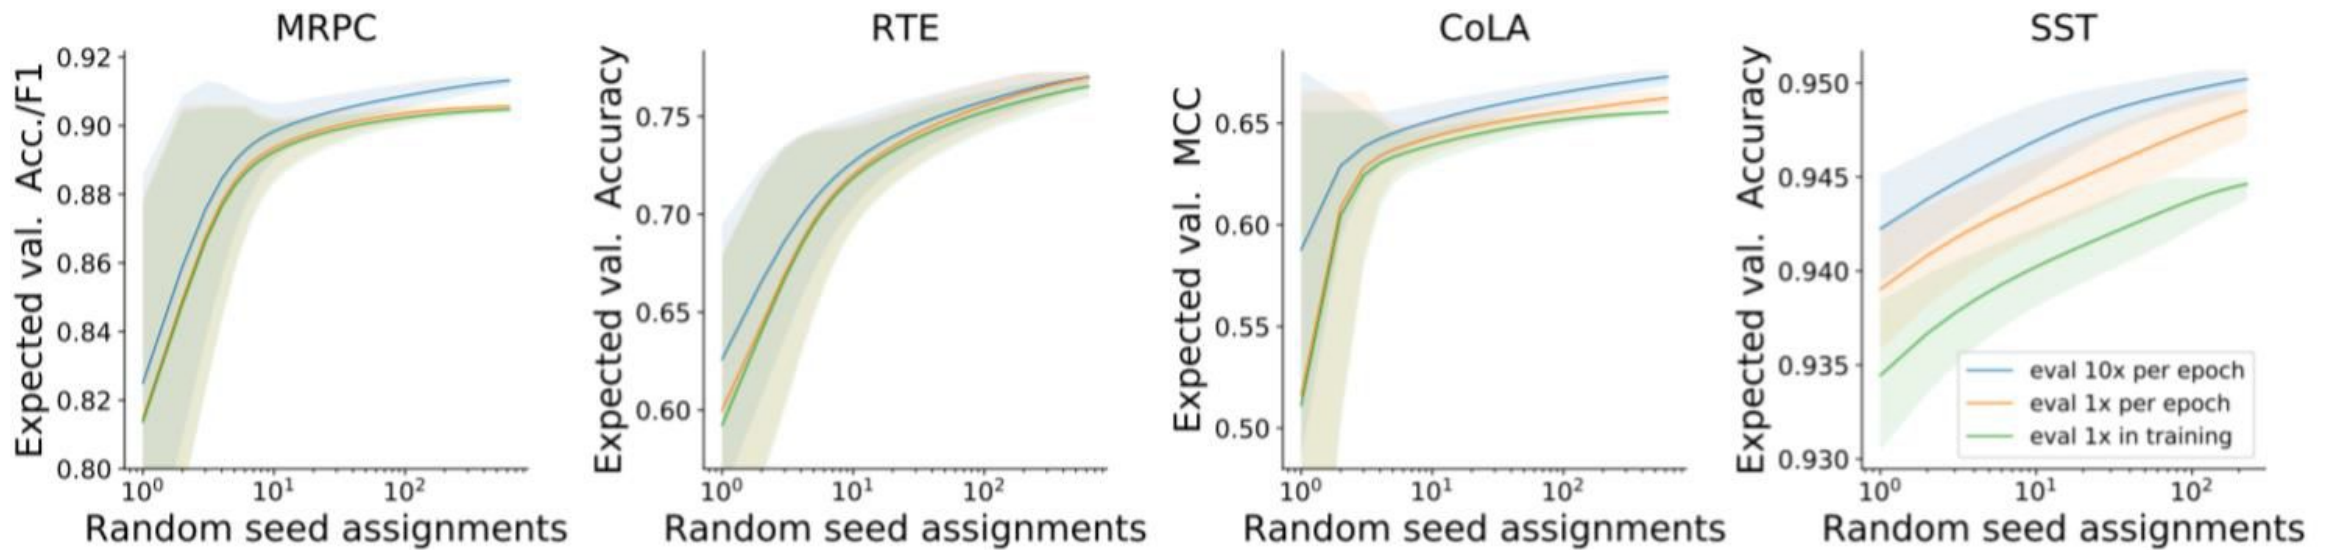
- There is a promising seeds
- These seeds can be distinguished early in training



Further Hyperparameters

3) Number of evaluation

- Expected validation performance as the number of evaluation increases

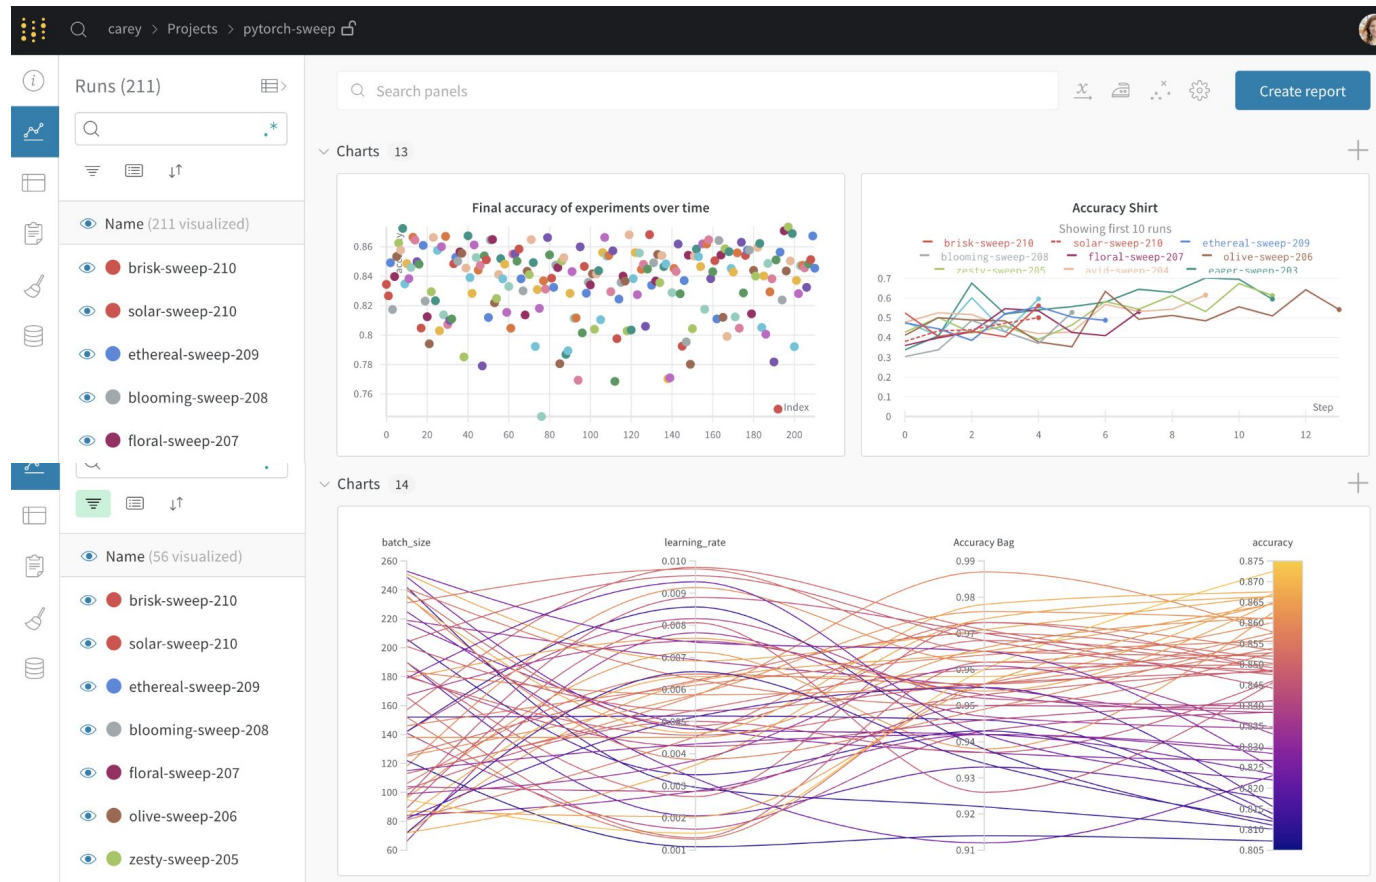


How to control randomness?

- `random.seed()`
- `np.random.seed()`
- `torch.manual_seed()`
- `torch.cuda.manual_seed()` / `torch.cuda.manual_seed_all()`
- `torch.backends.cudnn.deterministic = True`
- `torch.backends.cudnn.benchmark = False`
- [`torch.use_deterministic_algorithms\(True\)`](#)
- If you use CUDA tensors, we need to set the environment variable `CUBLAS_WORKSPACE_CONFIG` according to [CUDA documentation](#)

Wandb

- Experiment Tracking: <https://docs.wandb.ai/quickstart>
- Hyperparameter tuning: <https://docs.wandb.ai/guides/sweeps>



References

- http://cs231n.stanford.edu/slides/2019/cs231n_2019_lecture08.pdf